

AIAP Normative Assurance Standard

*Conformance requirements for claim-level assessment assurance in AI-mediated
higher education*

AIAP Release 6.5 | August 2026

Version-matched to the AIAP Core Working Paper v6.5 (Canonical Edition).

Contents

Status and use boundary.....	3
1. Scope.....	3
2. Controlling principles.....	3
3. Required roles and independence.....	3
4. Six-stage protocol.....	3
Stage 0 — Claim decomposition.....	3
Stage 1 — Learning claim.....	3
Stage 2 — Evidence-producing condition.....	4
Stage 3 — Assurance route.....	4
Stage 4 — Roles and evidence.....	4
Stage 5 — Programme mapping.....	4
5. Lane requirements.....	4
5.1 Lane 1 — Secured.....	4
5.2 Lane 2A — Competence-authenticated open.....	4
5.3 Lane 2B — Open non-certifying.....	4
5.4 Lane 3 — AI-integrated open.....	4
6. Evidence and authentication.....	5
6.1 Prompt security, exposure, and form equivalence.....	5
6.2 Evidence and authentication.....	5
7. Product-authentication combination.....	5
8. Dependence and independence.....	6
9. Modality and threat model.....	6
10. Programme and public claims.....	6
11. Capacity, accessibility, and appeals.....	6
12. Evidence governance.....	7
13. Temporal requalification.....	7
14. Validation boundary.....	7
15. Conformance statuses.....	7
Appendix A. Required record fields.....	8
Appendix B. Core-to-standard crosswalk.....	8

Status and use boundary

This standard is the normative companion to the AIAP Core Working Paper. It specifies what an institution **shall**, **should**, and **may** do when it claims conformance with AIAP Release 6.5. It does not establish empirical validity, confer legal authority, replace professional accreditation, or certify a programme merely because documents exist.

- **SHALL** indicates a requirement for the claimed scope.
- **SHOULD** indicates a strong default that may be varied only with a recorded rationale and equivalent safeguard.
- **MAY** indicates a permitted option.

1. Scope

The standard applies to consequential assessment claims about individual or programme capability where generative or agentic AI may affect production, evidence, authentication, or public interpretation. It governs the validity chain from learning claim to programme or public use.

2. Controlling principles

1. **Evidence-plan authority.** A lane, role, tier, or modality distinction **SHALL** have authority only when it changes the evidence plan or claim status needed for the permitted inference.
2. **Claim boundedness.** A claim **SHALL** identify the student or population, capability, scope, version, stakes, evidence-producing condition, evidence cut-off, decision rule, and permitted use.
3. **Non-transfer.** Evidence for artifact quality, present competence, production provenance, and AI-integrated capability **SHALL NOT** be treated as interchangeable.
4. **Fail-closed prerequisites.** A missing rights, identity, independence, construct-validity, accessibility, security, governance, or authority prerequisite **SHALL NOT** be compensated by stronger evidence elsewhere.
5. **Psychometric accumulation.** Genuinely informative and sufficiently independent observations **MAY** increase precision and generalizability when their dependence and construct coverage are declared.
6. **Public-claim integrity.** Programme, degree, accreditation, transcript, employer-facing, and public wording **SHALL NOT** exceed the evidence collected.
7. **Corrigibility.** Assurance **SHALL** expire, narrow, or be requalified after material changes or demonstrated failure.

3. Required roles and independence

The institution **SHALL** assign accountable owners for programme assurance, assessment design, delivery, authentication, moderation, accessibility, data governance, appeal, and public wording. Conflicts **SHALL** be disclosed. For high-consequence claims, the design **SHALL** include a materially independent review or evidence path appropriate to the threat model.

4. Six-stage protocol

Stage 0 — Claim decomposition

Every graded, progression-bearing, accreditation-bearing, or otherwise consequential material claim **SHALL** be identified separately. A whole-task lane **MAY** be used only when every material claim shares the same assurance logic.

Stage 1 — Learning claim

The capability and intended inference **SHALL** be explicit. The construct **SHALL** distinguish four positive assurance objects where relevant: artifact quality, present demonstrable competence, production provenance or unaided authorship, and AI-integrated capability. Lane 1 independent capability **SHALL** be recorded as a condition-qualified form of present demonstrable competence; Lane 2B non-certifying learning support **SHALL** be recorded as a claim-status routing value rather than as an additional positive assurance object.

Stage 2 — Evidence-producing condition

The condition SHALL be recorded as controlled or open, with modality, assistance channels, identity controls, residual threats, accessibility arrangements, and limitations. “Controlled” SHALL NOT be treated as perfect observation.

Stage 3 — Assurance route

Each material claim SHALL be assigned Lane 1, Lane 2A, Lane 2B, or Lane 3.

Stage 4 — Roles and evidence

AI roles, required evidence, authentication, combination model, dependence groups, retention, appeal, and reassessment SHALL be predeclared.

Stage 5 — Programme mapping

The result SHALL be mapped to module, programme, progression, award, accreditation, and public-claim uses. Under-evidenced claims SHALL be marked and corrected, narrowed, relocated, or withheld.

5. Lane requirements

5.1 Lane 1 — Secured

Lane 1 SHALL use performance under materially constrained conditions when an independent-capability inference is required. Identity, construct validity, accommodations, residual assistance, scoring reliability, and control strength SHALL be recorded. Lane 1 SHALL NOT be represented as absolute proof of unaided production or complete historical provenance.

5.2 Lane 2A — Competence-authenticated open

Lane 2A SHALL pair open work with claim-targeted competence authentication. Present demonstrable competence SHALL be treated as a declared profile rather than a presumed single trait. The programme SHALL name which of **explain, defend, verify, adapt, and transfer** are load-bearing for the claim. For consequential use, authentication SHALL include at least one materially unseen or meaningfully perturbed element. Rehearsal MAY support learning but SHALL NOT be the sole evidence of transfer. Lane 2A supports the declared competence profile and SHALL NOT be represented as complete production provenance.

Open-remote Lane 2A SHALL NOT be the sole warrant for a high-consequence claim unless a modality-specific threat model and controls establish sufficient environmental independence. Otherwise the institution SHALL add an independent anchor, reduce stakes, narrow wording, or reclassify the route.

5.3 Lane 2B — Open non-certifying

Lane 2B SHALL state explicitly that the task does not carry primary certification authority for the specified capability. The programme map SHALL identify where the consequential claim is assured elsewhere. Repeated, heavily weighted, progression-bearing, or programme-critical use SHALL trigger promotion to Lane 1, 2A, or 3, or removal of the unsupported claim.

5.4 Lane 3 — AI-integrated open

Lane 3 SHALL make accountable AI-mediated practice part of the construct. Evidence SHALL identify delegated functions, retained human decisions, verification, rejected recommendations, override/stop points, risk controls, and responsibility. For high-consequence claims of judgement or

oversight, workflow records alone are insufficient; an authenticated judgement or transfer component SHALL be included.

6. Evidence and authentication

6.1 Prompt security, exposure, and form equivalence

For every consequential Lane 2A event, the institution SHALL declare the cohort and exposure window, prompt or live-perturbation source, allocation method, use count, rotation or retirement rule, exposure-monitoring route, and responsible owner. The unseen element SHOULD be generated through a claim-targeted live perturbation of the student's own artifact where that preserves construct coverage and accessibility. Static banks MAY be used only under governed access and exposure controls.

Different forms carrying the same hurdle decision SHALL have a common construct blueprint and evidence supporting parallel-form equivalence. Where equivalence is not established, the institution SHALL narrow the inference, use a common equivalent challenge, or treat form as a material source of uncertainty. Credible circulation or compromise SHALL trigger item retirement, incident review, and requalification of affected decisions.

The interval between submission and authentication SHALL be recorded and justified. A delay that materially changes the relationship between the artifact and present competence SHALL narrow the inference or require a new evidence event.

6.2 Evidence and authentication

The evidence plan SHALL distinguish product, process, source, reflection, identity, competence, provenance, transfer, and judgement evidence. An AI-use record SHALL be required only when it materially serves a declared learning, validity, accountability, or governance purpose. Data SHALL NOT be collected merely because it is available.

Weak authentication SHALL be routed first as a validity concern. Misconduct referral requires independent evidence of deception, fabrication, outsourcing contrary to a declared production constraint, identity breach, or violation of controlled conditions.

For collaborative work, a group artifact SHALL NOT carry an individual consequential claim unless each student's evidence is separately attributable. Group walkthroughs MAY reduce scheduling burden only when each student receives a separately addressed challenge, is separately scored, and has an individual decision record. Where an individual consequential claim depends on authentication, every student SHALL complete the declared minimum authentication; sampling SHALL NOT substitute for individual certification evidence and MAY be used only for moderation, calibration, audit, or research.

7. Product-authentication combination

The institution SHALL select and publish one of the following before assessment begins:

- **Hurdle model:** certification requires a minimum authentication result; product strength cannot compensate for failure of the hurdle.
- **Dual-score model:** product quality and authenticated competence are scored and reported separately, with a declared progression rule.
- **Integrated rubric model:** evidence sources are mapped to criteria, double counting is prohibited, and essential minima are declared.

Thresholds SHALL be justified through a local standard-setting argument, calibration, consequence analysis, equivalent reassessment, and appeal. For Lane 2B, no product-authentication combination model applies because the corresponding capability claim is non-certifying in that task; the machine-readable record encodes this state as non_certifying.

8. Dependence and independence

Every load-bearing evidence object SHALL record source, method, time, claim scope, directness, independence group, uncertainty, accessibility route, retention, and challenge status. Source independence and environmental independence SHALL be assessed separately. Contradictory evidence SHALL be preserved and adjudicated rather than averaged away.

9. Modality and threat model

The threat model SHALL identify plausible production delegation, identity substitution, covert assistance, rehearsed response, process-log fabrication, model drift, accessibility conflicts, and data-governance risk. Controls SHALL be the least intrusive sufficient means for the claim. Remote proctoring or surveillance SHALL NOT be adopted merely because it is technically available.

10. Programme and public claims

The programme SHALL maintain a versioned assurance map showing where each programme learning outcome is secured, competence-authenticated, AI-integrated, non-certifying, or under-evidenced. Broad claims SHALL be supported across sufficient prompts, cases, occasions, assessors, and modalities. One brief event SHALL NOT be generalized beyond its demonstrated scope without evidence addressing case/content specificity.

The programme SHALL publish or make inspectable an **assurance composition record**: the proportion of mapped credit or decision weight whose person-attributable warrant rests on materially constrained performance, present-competence authentication, AI-integrated capability, and non-certifying activity. The mix SHALL be justified against disciplinary, professional, accreditation, and public-claim expectations. A larger percentage in one route SHALL NOT be represented as higher quality by itself.

11. Capacity, accessibility, and appeals

A consequential claim SHALL pass the Assurance Capacity Gate. The thirteen canonical conditions are:

1. G01 — Claim map approved: material claims, stakes, conditions, lane, certification object or routing status, programme use, and public wording are recorded.
2. G02 — Prompt security governed: the bank or live-perturbation method, cohort, exposure window, allocation, use counts, rotation, retirement, monitoring, and ownership are declared.
3. G03 — Parallel-form equivalence supported: forms share a construct blueprint and carry evidence sufficient for the declared decision, or non-equivalence is explicitly bounded.
4. G04 — Submission-authentication interval justified: the interval is recorded and remains defensible for the present-competence inference.
5. G05 — Assessor qualification and calibration established: assessors have relevant expertise, rubric calibration, accessibility and integrity competence, and fatigue/drift controls.
6. G06 — Reviewer and panel capacity sufficient: required hours, throughput, reserve margin, moderation, and calendar fit are feasible.
7. G07 — Independence and conflicts controlled: responsibilities, conflicts, recusal, and an independent review or evidence path are operational.
8. G08 — Accessibility-equivalent route operational: an equivalent route, construct map, scheduling, notice, and residual differences are governed.
9. G09 — Appeal and second-review route operational: rubric and rationale access, lawful evidence access, non-involved review, and rehearing or equivalent reassessment are available.
10. G10 — Data governance approved: purpose, minimum necessary data, access, retention, correction, disposal, and incident response are authorized.
11. G11 — Capacity-degradation fallback declared: mid-cycle illness, rescheduling, technology, prompt, recording, accessibility, and data failures have governed responses.
12. G12 — Threat model acceptable: modality-specific assistance, substitution, rehearsal, process-log, accessibility, privacy, and agentic-AI threats are addressed by proportionate controls.

13. G13 — No unresolved high-severity exception: no remaining threat, rights, validity, accessibility, governance, or authority exception defeats the claim.

A failed condition blocks the claim or requires redesign, reduced stakes, narrower scope, delay, or reclassification. Resource scarcity SHALL NOT lower a rights, validity, or accessibility floor.

If capacity degrades after delivery begins - including assessor illness, mass rescheduling, technology failure, prompt compromise, recording loss, or accessibility failure - the institution SHALL pause the affected assurance route, use a predeclared construct-equivalent fallback where available, or delay, narrow, or withhold the claim. It SHALL NOT silently weaken the minimum evidence to preserve a timetable.

12. Evidence governance

The institution SHALL document purpose, lawful or institutional basis, minimum necessary data, access roles, security, retention, correction, appeal, deletion/disposal, incident response, and an equivalent non-surveillance route where needed. Disability or accommodation information SHALL receive appropriate protected handling. Local legal review remains necessary. Prompt banks, exposure logs, form blueprints, recordings, assessor notes, and appeal materials SHALL be protected according to their security, privacy, and evidentiary function.

A student challenging a consequential authentication decision SHALL have access, subject to lawful and proportionate limits, to the applicable rubric, decision rationale, recording or assessor notes where they exist, and the route to review by an assessor not involved in the original decision. A rehearing or equivalent reassessment SHALL be available where procedural unfairness, technology failure, accessibility failure, or prompt compromise is established.

13. Temporal requalification

Every current assurance claim SHALL identify protocol version, assessment version, scope, evidence set, authority, exceptions, review date, and triggers. Requalification is required after material changes to learning outcomes, task design, tool/model environment, modality, population, evidence method, policy, or observed failure.

14. Validation boundary

An institution MAY claim implementation or conformance only for the scope actually evidenced. It SHALL NOT claim that AIAP itself is empirically validated merely because the standard is followed. Pilot, validation, null, adverse, burden, and withdrawal findings SHALL remain publishable and capable of changing the protocol.

15. Conformance statuses

- **CONFORMS_FOR_DECLARED_SCOPE** — every applicable SHALL requirement is evidenced for the exact scope.
- **CONFORMS_WITH_VERIFIED_CONTROLS** — conformance depends on named, owned, monitored, time-bounded controls.
- **PARTIAL_IMPLEMENTATION** — selected elements are used, but a full conformance claim is not permitted.
- **INDETERMINATE_INSUFFICIENT_EVIDENCE** — adequacy or defect cannot be established safely.
- **NONCONFORMING_PATCHABLE** — a material, repairable defect exists.
- **NONCONFORMING_FOR_STATED_CLAIM** — the stated claim conflicts with a material requirement.
- **REFUSAL_TO_ASSESS_CONFORMANCE** — competence, independence, access, capacity, version, or governance prerequisites are invalid.

Appendix A. Required record fields

Record	Minimum fields
Claim profile	claim ID, capability, stakes, condition, lane, modality, evidence, combination model, permitted inference, use
Authentication plan	construct profile; load-bearing verbs; prompt/live-perturbation source; cohort and exposure window; allocation, rotation, retirement and exposure monitoring; form-equivalence evidence; submission-authentication interval; identity controls; scoring; moderation; accessibility; reassessment; appeal
Evidence object	source, method, date/time, claim scope, directness, independence group, uncertainty, retention, challenge status
Programme map	PLO, evidence points, lane, weight, assurance-composition share, progression relevance, redundancy, under-evidenced status, public wording
Capacity record	inputs, assumptions, thirteen gate conditions, degradation fallbacks, controls, owner, expiry
Requalification record	trigger, review date, authority, evidence, disposition, superseded version

Appendix B. Core-to-standard crosswalk

The version-matched crosswalk in AIAP_v6.5_CORE_TO_STANDARD_CROSSWALK.md identifies the controlling Core-paper section, Standard requirement, implementation artifact, machine-readable field where applicable, and verification evidence for each principal rule.