

AIAP Implementation Handbook

From claim decomposition to programme assurance

AIAP Release 6.5 | August 2026

Version-matched to the AIAP Core Working Paper v6.5 (Canonical Edition).

Contents

How to use this handbook.....	3
1. AIAP-Lite: minimum viable adoption.....	3
2. AIAP-Full: programme governance.....	3
3. Assessment redesign workflow.....	3
Step 1 — Diagnose the current inference.....	3
Step 2 — Run the failure test.....	3
Step 3 — Choose the repair.....	3
Step 4 — Select the combination model.....	3
Step 5 — Test modality.....	4
Step 6 — Pass the capacity gate.....	4
4. Worked failure and repair.....	4
5. Lane design cards.....	4
Lane 1.....	4
Lane 2A.....	4
Lane 2B.....	4
Lane 3.....	4
6. Authentication design.....	5
6.1 Prompt security and form equivalence.....	5
7. Programme assurance.....	5
8. Online-only provision.....	5
9. Workload and scheduling.....	5
10. Quality-assurance cycle.....	6
10.1 Group and collaborative assessment.....	6
10.2 Minimum appeal record.....	6
11. Minimum implementation records.....	6
Appendix A. Implementation checklist.....	6

How to use this handbook

The handbook operationalizes the Core Working Paper and Normative Assurance Standard. It is not an independent source of authority. Start with AIAP-Lite for one bounded assessment; use AIAP-Full when claims affect progression, awards, accreditation, safety, or public trust.

1. AIAP-Lite: minimum viable adoption

1. List every material learning claim.
2. State what each claim is intended to certify.
3. Record whether evidence is produced under controlled or open conditions.
4. Select Lane 1, 2A, 2B, or 3 for each claim.
5. Specify the minimum evidence and any authentication.
6. State what the result may and may not support.
7. Give students a short, plain-language instruction.
8. Check feasibility, accessibility, reassessment, and appeal before release.

A whole-task label is acceptable only when all material claims use the same logic.

2. AIAP-Full: programme governance

AIAP-Full adds a programme assurance map, evidence-dependence groups, authentication-event design, workload modelling, modality threat models, disclosure-bias checks, annual requalification, external wording control, and a pilot/validation plan.

3. Assessment redesign workflow

Step 1 — Diagnose the current inference

Ask: what are we currently inferring from the submitted product? Which parts depend on unobserved compliance? Which claims are independent, open, AI-integrated, or merely developmental?

Step 2 — Run the failure test

A task fails when a consequential capability claim rests only on an open product, a declaration, an unenforceable boundary, or a record that can be reconstructed without the student possessing the claimed capability.

Step 3 — Choose the repair

- Add materially constrained performance → Lane 1.
- Add claim-targeted competence authentication → Lane 2A.
- Remove certification authority and map the claim elsewhere → Lane 2B.
- Make AI-mediated judgement, verification, and control part of the construct → Lane 3.

Step 4 — Select the combination model

Choose hurdle, dual-score, or integrated rubric for certifying claims. Publish thresholds, non-compensation rules, reassessment, appeal, and integrity routing. For Lane 2B, no combination model applies because the capability is explicitly non-certifying in that task; record the machine status as non_certifying and map the capability elsewhere.

Step 5 — Test modality

In-person, invigilated-remote, and open-remote events are not equivalent. Record source and environmental independence separately. Include at least one materially unseen or perturbed element for consequential Lane 2A.

Step 6 — Pass the capacity gate

Use the version-matched workbook. A false green occurs when a favourable minutes-per-student estimate masks calendar, staffing, accessibility, independence, or appeal failure.

4. Worked failure and repair

Initial task: A 2,500-word strategy report says AI may be used for brainstorming but not drafting. The task is completed off campus, and no process or competence evidence is required.

Defect: The distinction is communicated but not structurally observable. The product can support output-quality judgement, but it cannot by itself support an independent-capability or production-provenance claim.

Repair A — Lane 2A: retain the open report, add a source-verification table and an authenticated unseen adaptation. Use a hurdle or dual-score model.

Repair B — Lane 3: assess AI-workflow design, verification, human correction, risk controls, and unseen audit of a new output.

Repair C — Lane 2B: use the report for feedback and practice only; assure the consequential capability in a separate secured or authenticated assessment.

5. Lane design cards

Lane 1

Use for essential independent knowledge, safety-critical performance, live procedural skill, or foundational capability. Record control strength and residual limitations. Avoid construct narrowing merely to increase security.

Lane 2A

Use for complex open work where present competence matters. Authentication should target the declared competence profile: explain, defend, verify, adapt, and transfer. Programmes name which components are load-bearing for the claim. Rehearsed preparation is legitimate; the unseen or perturbed element carries the transfer burden.

Lane 2B

Use for formative, remediable, or triangulated activity. Put the non-certifying status in the brief and programme map. Review accumulated credit weight and programme reliance each year.

Lane 3

Use when responsible AI-mediated practice is part of the learning outcome. Require workflow, verification, retained decision rights, rejected recommendations, accountability, and—where consequential—authenticated judgement or transfer.

6. Authentication design

A strong event has a declared competence profile, a protected prompt or live-perturbation design, an unseen/perturbed component, analytic scoring guide, assessor qualification and calibration, moderation and drift checks, an equivalent accessible route, reassessment, appeal, data-minimization plan, and modality threat model.

6.1 Prompt security and form equivalence

1. Declare the material claim, load-bearing competence component, cohort, scheduling window, and likely exposure path.
2. Prefer artifact-specific live perturbations - changed constraints, new data, planted errors, or changed cases - where they preserve the construct and accessibility.
3. If a bank is used, record prompt ID, form ID, allocation method, use count, access owner, rotation and retirement rule, and exposure status.
4. Monitor circulation through student reports, statistical anomalies, staff observations, and public leakage. Retire compromised prompts and review affected decisions.
5. Establish parallel-form equivalence through a common blueprint, expert review, pilot evidence, decision-consistency evidence, and accessibility comparison. If equivalence is not supported, do not use the forms interchangeably for the same hurdle.
6. Record and justify the interval between submission and authentication. Long delays may change what the event can say about the artifact.

No universal bank-size ratio is imposed. The programme must instead show why its bank or live-perturbation method remains adequate for the cohort, exposure window, stakes, and threat model.

Weak performance is not proof of misconduct. Route first to the grading/validity process; refer to integrity only when separate evidence indicates deception, fabrication, identity breach, prohibited outsourcing, or controlled-condition violation.

7. Programme assurance

Map each programme learning outcome across years and evidence types. Avoid concentrating all assurance in one exit event. Broad claims require distributed evidence across prompts/cases, occasions, assessors, and contexts. Mark every Lane 2B contribution and every under-evidenced claim. Publish or retain for audit the assurance-composition ratio: the share of mapped credit or decision weight carried by Lane 1, Lane 2A, Lane 3, and Lane 2B. Justify the mix; do not treat a route's percentage as a quality score.

8. Online-only provision

Online provision can support consequential claims only when it supplies a sufficiently independent evidence path for the stated construct. Options include authorized local test centres, validated invigilated-remote assessment, observed workplace/practicum evidence, distributed supervised checkpoints, and narrower public wording. Where no adequate path exists, the programme must narrow the claim rather than describe open-remote evidence as equivalent to in-person control.

9. Workload and scheduling

Model staff-hours, elapsed calendar time, throughput, moderation, training, accessibility, rescheduling, recording review, appeals, and data governance separately. Include transition time, maximum continuous session length, scheduled breaks, assessor qualification, and drift checks. Compare the proposed route with the assessment it replaces or supplements rather than treating the authentication cost in isolation. Parallel staffing changes elapsed time, not total staff-hours. Record student preparation, waiting, travel/setup, accessibility, rescheduling, and anxiety burden in pilots.

10. Quality-assurance cycle

Before delivery: claim map, threat model, student comprehension test, calibration, capacity gate, privacy/accessibility review.

During delivery: incident log, prompt-exposure monitoring, accommodation monitoring, assessor fatigue and drift checks, rescheduling, data minimization. If capacity fails mid-cycle, use only a predeclared construct-equivalent fallback; otherwise delay, narrow, or withhold the affected claim.

After delivery: reliability and decision review, mismatch analysis, disclosure bias, appeals, burden, public wording, requalification decision.

10.1 Group and collaborative assessment

A shared artifact may support a claim about the group's output or collaborative process, but it cannot by itself certify each student's capability. Decompose the individual and group claims separately. A group walkthrough is acceptable only when every student receives a separately addressed challenge, is separately scored, and has an individual decision record. Collaborative process evidence may remain Lane 2B when it supports learning but does not carry individual certification.

10.2 Minimum appeal record

Students should receive the applicable rubric, decision rationale, and access to the recording or assessor notes where these exist and may be shared lawfully. Review should be conducted by a qualified person not involved in the original decision. A rehearing or equivalent reassessment is required where procedural unfairness, technology failure, accessibility failure, or prompt compromise is established.

11. Minimum implementation records

Use the companion toolkit, workbooks, schema, and checklists. No local form may silently change the four lanes, claim boundaries, non-compensation rules, or public-claim discipline.

Appendix A. Implementation checklist

- ☐ Material claims decomposed.
- ☐ Learning claim, positive assurance object, and any routing status stated.
- ☐ Condition and modality recorded.
- ☐ Lane assigned for each claim.
- ☐ Evidence and combination model predeclared.
- ☐ Unseen/perturbed component included where consequential Lane 2A applies.
- ☐ Prompt exposure window, allocation, rotation, retirement, and form-equivalence evidence recorded.
- ☐ Submission-authentication interval recorded and justified.
- ☐ Group work produces separately attributable individual evidence where individual claims are made.
- ☐ Assessor qualification, session limits, breaks, and drift checks planned.
- ☐ Mid-cycle capacity-degradation fallback declared.
- ☐ Accessible equivalent route designed.
- ☐ Capacity Gate passed.
- ☐ Student wording comprehension-tested.
- ☐ Programme map and external wording updated.
- ☐ Requalification trigger and review date set.